



Regularization

Justin Ponders

Objectives

- Define and describe regularization for regression models
- Write the regularized loss function
- Describe how regularization affects regression coefficients
- Describe the differences between the Lasso, Ridge, and ElasticNet models
- Implement and visualize the penalties using sklearn

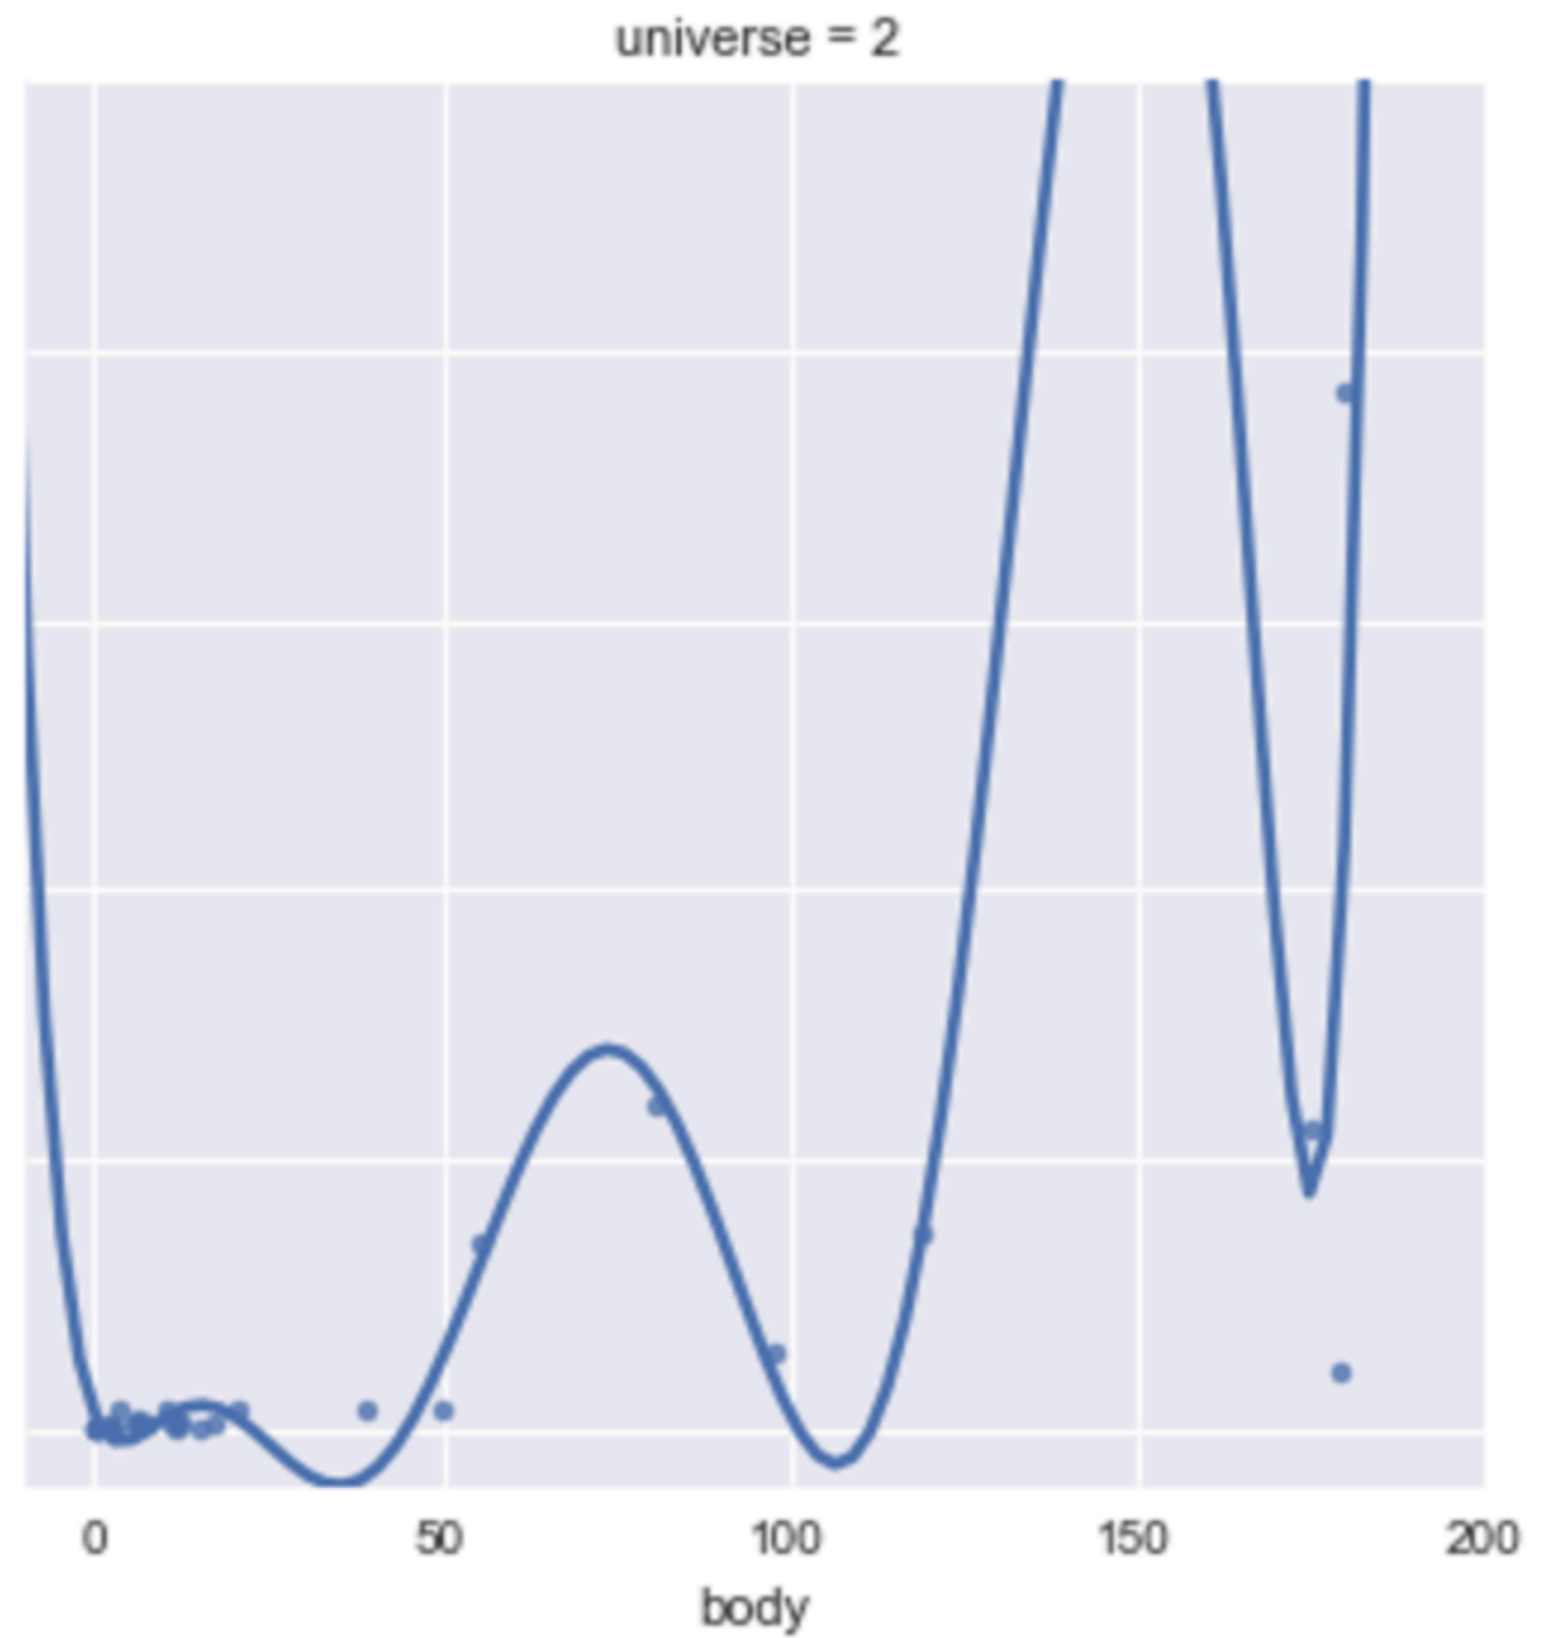
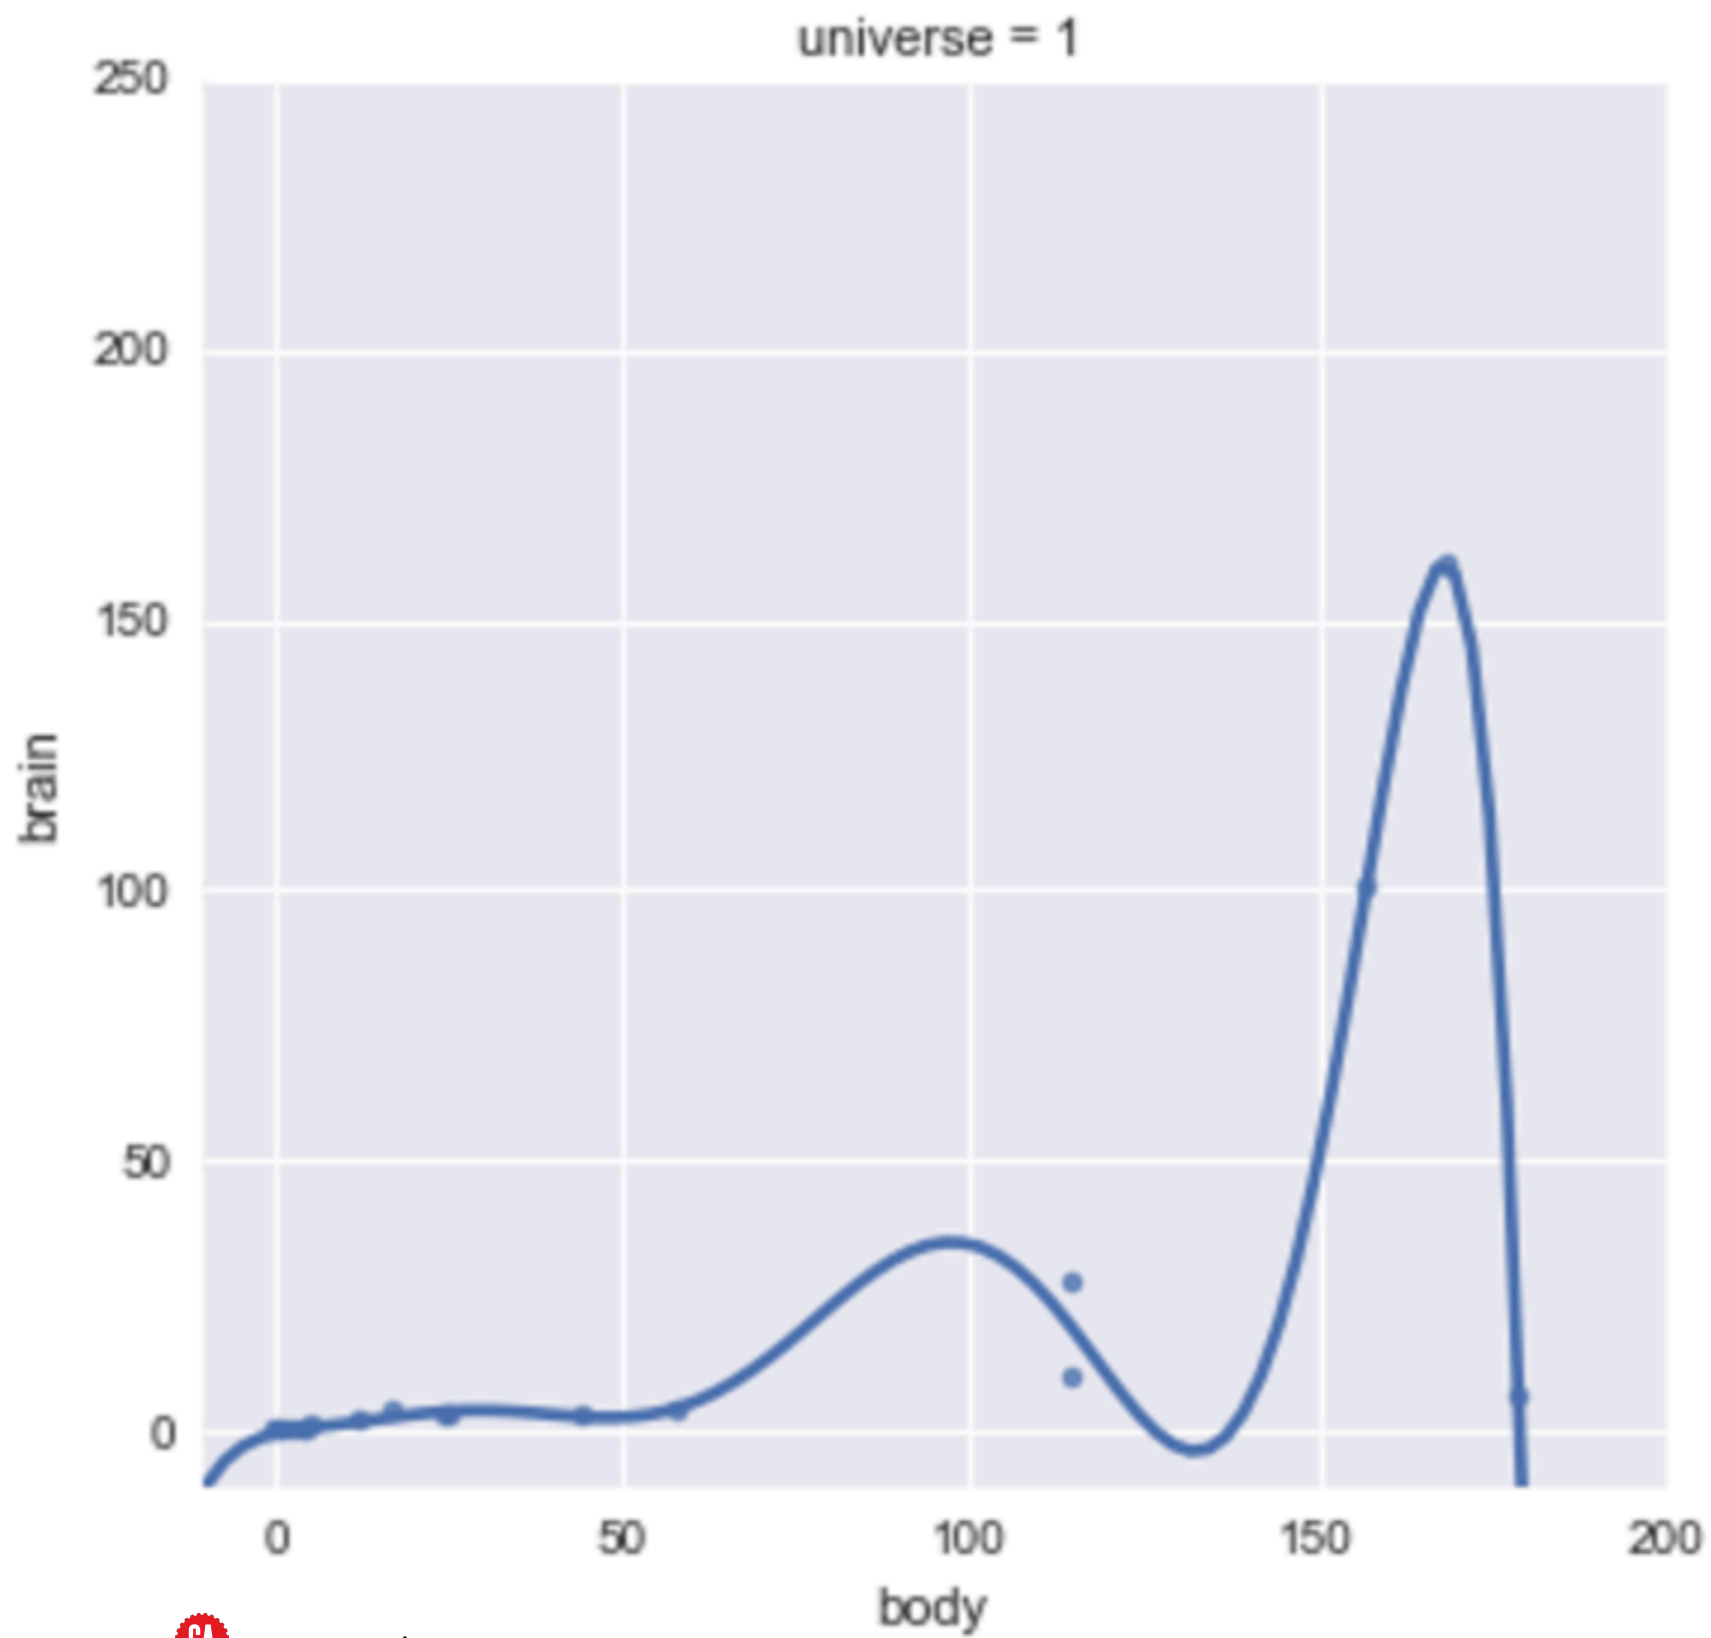
Warmup

With a partner, describe what overfitting is and how it occurs. What is the impact of overfitting?

Overfitting

Overfitting means building a model that matches the training data "too closely." The model ends up training on noise rather than signal.

- Usually cause by model that is too complex
- Overfit model does not generalize
- Low bias/high variance models



Do I need to worry about overfitting with Linear Regression?

"Good" properties

- Low complexity
- High bias/low variance
- Does not tend to overfit

Do I need to worry about overfitting with Linear Regression?

Danger zone

- Including irrelevant features (signal v noise)
- p (number of features) is close to n (number of observations)
- Correlated inputs
- Numerically large coefficients

What
Do I Do
Now

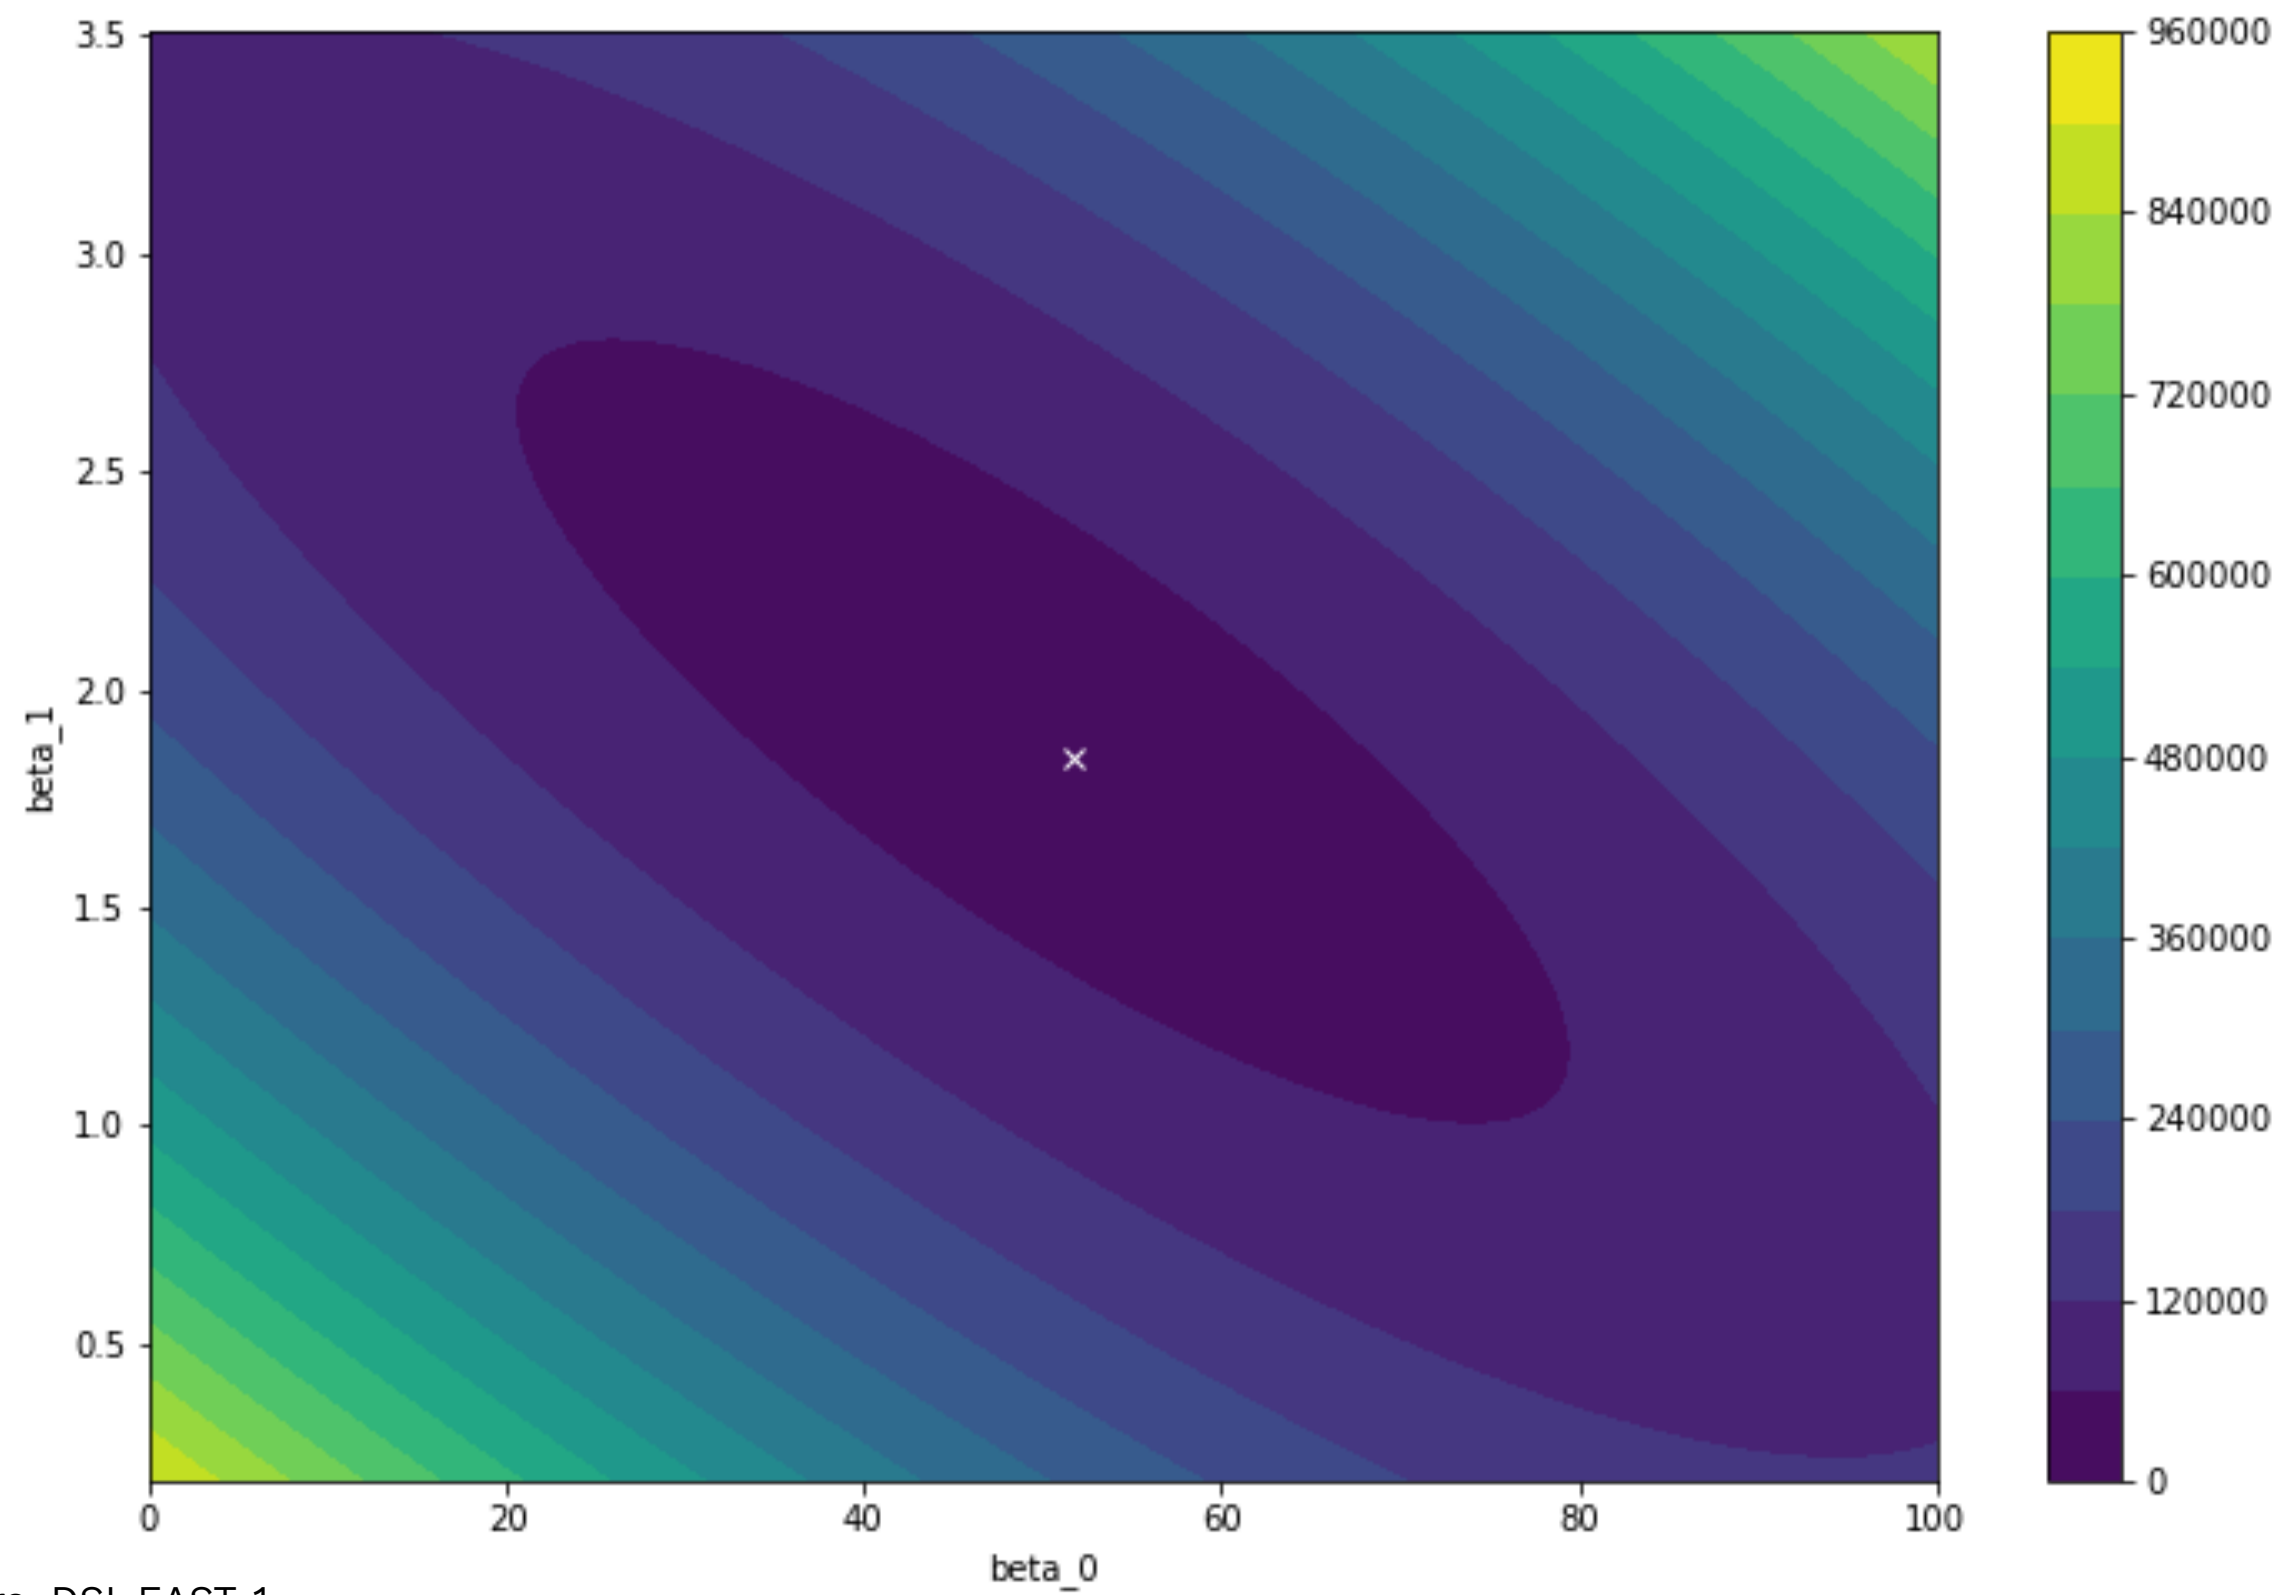


Error and the Loss Function

$$\begin{aligned} RSS(\beta_0, \beta_1) &= \sum_{i=1}^n (y_i - \hat{y}_i)^2 \\ &= \sum_{i=1}^n (\hat{y}_i - \beta_0 - \beta_1 x_i)^2 \end{aligned}$$

The goal of training is to minimize $RSS(\beta_0, \beta_1)$, i.e.

$$\beta_0, \beta_1 = \arg \min RSS(\beta_0, \beta_1)$$



Ridge Regression

(a.k.a. Tikhonov regularization, weight decay, L_2 regularization)

$$J(\beta_0, \beta_1) = RSS(\beta_0, \beta_1) + \alpha\beta_1^2$$

Ridge regression **penalizes** the model for having large coefficients. As α increases, β_1 will decay.

α acts as a "tuning" parameter.

Ridge Regression (general case)

$$J(\beta_0, \beta_1, \dots, \beta_p) = RSS(\beta_0, \beta_1, \dots, \beta_p) + \alpha \sum_{i=1}^p \beta_i^2$$

Ridge **shrinks** the regression coefficients.

Ridge Regression

Get ready to roll (down the loss function!)

Check: Find the sklearn documentation on Ridge Regression. Locate the **model description** (inputs, outputs, parameters, methods), and the **discussion of the theory** with examples.

1. How do you set the regularization strength?
2. How can you get the values of the regularized regression coefficients?

Lasso Regression

a.k.a. L_1 regularization

$$J(\beta_0, \beta_1) = RSS(\beta_0, \beta_1) + \alpha |\beta_1|$$

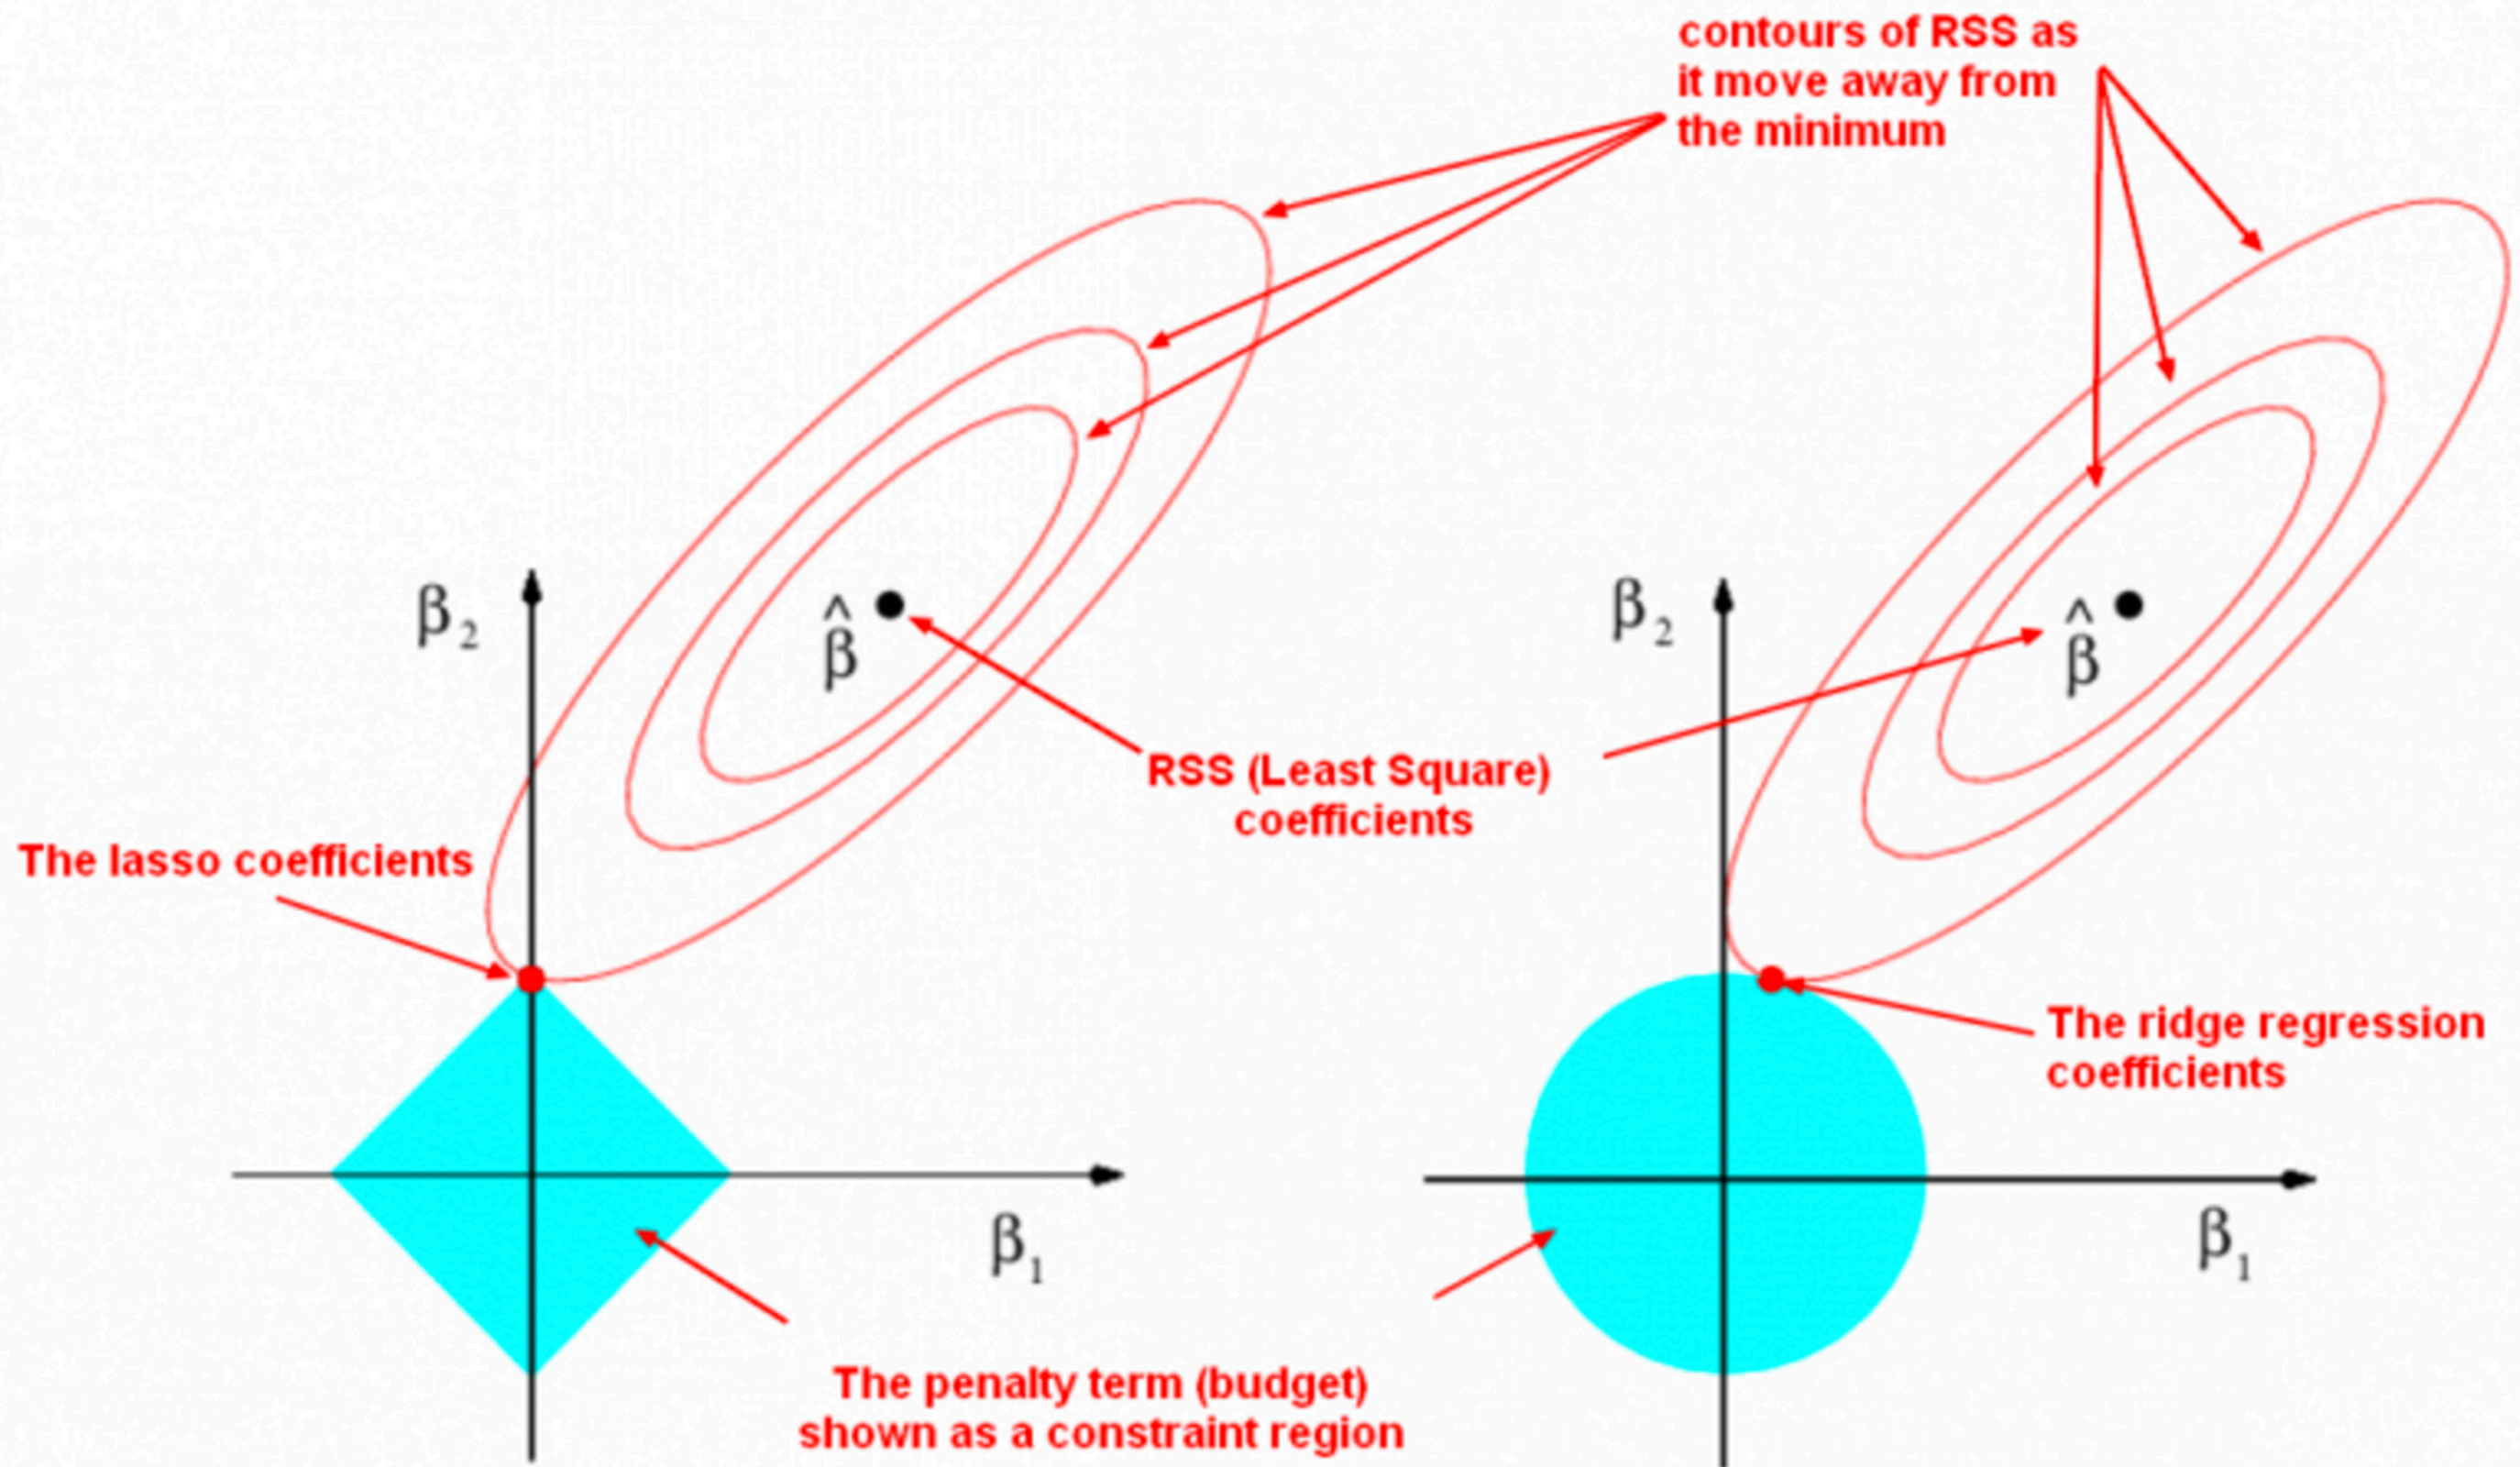
Lasso regression **penalizes** the model for having large coefficients. As α increases, β_1 will decrease, even to the point of zero.

α acts as a "tuning" parameter.

Lasso Regression (general case)

$$J(\beta_0, \beta_1, \dots, \beta_p) = RSS(\beta_0, \beta_1, \dots, \beta_p) + \alpha \sum_{i=1}^p |\beta_i|$$

Ridge *shrinks* the regression coefficients, and may "zero-out" unimportant features.



Tuning Bias vs Variance with Regularization

Alert!! Key takeaway!

- Increase α (turn **up** regularization)
 - Increase bias
 - Decrease variance
- Decrease α (turn **down** regularization)
 - Decrease bias
 - Increase variance

Additional Considerations

- Features (inputs) should be **standardized** in regularized models
 - **Why?**
- Ridge vs Lasso?
 - Maybe have irrelevant features? **Lasso**
 - Just want the best prediction? **Try both**
 - Want to use both? **ElasticNet**

Elastic Net Regression

$$J(\beta_0, \beta_1) = RSS(\beta_0, \beta_1) + \alpha_1 |\beta_1| + \alpha_2 \beta_1^2$$

Elastic net combines Ridge and Lasso penalties

α_1 and α_2 both act as "tuning" parameters

Elastic Net Regression

A second (equivalent formulation) used by sklearn

$$J(\beta_0, \beta_1) = RSS(\beta_0, \beta_1) + \alpha\rho |\beta_1| + \frac{\alpha(1 - \rho)}{2} \beta_1^2$$

Elastic net combines Ridge and Lasso penalties

- α = penalty strength
- ρ = Lasso (L_1) ratio

Elastic Net Regression

A second (equivalent formulation) used by sklearn

$$J(\beta_0, \beta_1) = RSS(\beta_0, \beta_1) + \alpha\rho |\beta_1| + \frac{\alpha(1-\rho)}{2} \beta_1^2$$

Elastic net combines Ridge and Lasso penalties

Check: what values of ρ lead to (a) Ridge and (b) Lasso regression?

Elastic Net Regression (general case)

$$J(\beta_0, \beta_1, \dots, \beta_p) = RSS(\beta_0, \beta_1, \dots, \beta_p) + \alpha\rho \sum_{i=1}^p |\beta_i| + \frac{\alpha(1-\rho)}{2} \sum_{i=1}^p \beta_i^2$$

[Elastic Net] allows for learning a sparse model where few of the weights are non-zero like Lasso, while still maintaining the regularization properties of Ridge.

— sklearn docs

Turn and Talk

With a partner, discuss and summarize...

1. What is Ridge regularization?
2. What is Lasso regularization?
3. What is Elastic Net regularization?

After 4 minutes, I will call on volunteers to summarize one of these responses.

Python Time